

User Manual

metabolopan — Metabolomic Enrichment Analysis

Contents

Stage 1 — Input parsing	3
MS-DIAL .txt	3
Group mapping .csv	3
Missing values (NaN) vs true zeros (0.0)	3
Stage 2 — Sample normalization	5
Stage 2 — Deduplication by InChIKey	7
Cascade decision table	7
Audit CSV	8
Opt-out	8
Stage 2 — DAM (Differentially Accumulated Metabolites)	8
Stage 3 — Enrichment (over-representation analysis)	11
Enrichment Analysis setup screen	11
Pathway mode	12
Module mode	14
Starting a new analysis round	15
Dual-mode (positive + negative ionization) input	15
When to use dual-mode	15
Preparing inputs	15
Unbalanced or missing-mode samples	16
Stage 1 UI	16
Stage 2 (shared setup, per-mode DAM)	17
Stage 3 — dual-mode N and K math	17
Worked example	17
Caches and provenance	18
Saving and loading session settings (reproducibility)	18
What's in the file	19
When is each button available	19
Loading workflow	19
What if I load settings before uploading metadata?	19
Hand-editing the JSON	19
Reporting bugs	20
Key references	20

This manual documents what the software does numerically — algorithms, default thresholds, deviations from common alternatives — so you can defend any number it produces in a paper or report. Read this once before publishing results that depend on the software.

The pipeline operates in three stages. Default thresholds appear in square brackets [...].

Stage 1 — Input parsing

metabolopan takes input in one of two modes, chosen by how many MS-DIAL .txt files you load. **Single-mode** is one MS-DIAL .txt + one group-mapping .csv; **dual-mode** is two MS-DIAL .txt files (one positive, one negative ionization) + one group-mapping .csv that includes a biosample column. The file formats below apply to both modes; the dual-mode-specific mechanics (biosample pairing, group parity checks, per-mode DAM) are covered in [Dual-mode \(positive + negative ionization\) input](#) below.

MS-DIAL .txt

- First 4 rows are MS-DIAL metadata (Class, File type, Injection order, Batch ID); the 5th row is the column header. A column is treated as a real sample injection — and kept in sample_cols — when its File type value is non-empty AND not "NA" AND not the literal row label "File type". This **includes** Sample and Blank (process blanks); it excludes only MS-DIAL's per-group Average / Stdev aggregation columns (labelled NA).
- **Version compatibility.** Both MS-DIAL 4 and MS-DIAL 5 Alignment exports are supported. Column lookup is by name, so MS-DIAL 5's reordered/renamed scoring columns (it splits Dot product into Simple / Weighted dot product) parse identically; metabolopan uses only the columns shared by both versions.
- **Missing values.** Empty / whitespace / "null" / "NA" / unparseable intensity cells become f64::NAN. Explicit "0" stays 0.0. This prevents downstream statistics from confusing missing measurements with true zeros.

Group mapping .csv

- The CSV must contain a column named sample and a column named group, in **any position/order**. An optional biosample column (any position; required for dual-mode — see [Dual-mode input](#) below) is recognized by name. Any further columns are parsed as optional metadata. Column names are matched exactly (case-sensitive); a missing sample/group column, a duplicated sample/group/biosample column, empty group cells, or duplicate sample names are each rejected with a descriptive error. Metadata columns are classified per-column at load time: a column whose non-empty cells all parse as numbers is exposed to Stage 2's "Metadata column" normalization ratio (e.g. dry_weight, dilution, total_protein); a column with any non-empty non-numeric cell (e.g. a biosample label like CTR-01) is silently dropped from that ratio and a WARN line in the in-app log pane names which column was skipped and how many cells failed to parse. Empty metadata cells parse to None. Samples that appear in the MS-DIAL .txt but not in the CSV are flagged Unassigned; rows in the CSV that name samples missing from the .txt are logged as warnings and ignored. **Unassigned samples are visible on Stage 1 only** — the input-summary panel shows them with a yellow Unassigned (N samples) row so you know they exist, but they are dropped from the working matrix when you click Start DAM on the Stage 2 setup screen. No normalization, deduplication, DAM statistic, or downstream export ever sees them. To include a sample in analysis, add it to the metadata CSV with a real group label; to exclude a sample entirely (so it doesn't even show on Stage 1), remove its column from the MS-DIAL .txt File type row (set the entry to NA). Metadata column order in the Stage 2 dropdown matches the CSV header order (not alphabetical), so the user sees columns in the order they wrote them.
- **Stage 1 → Stage 2 gate.** The Continue to DAM button stays disabled until: both files parse successfully, the slot-#1 ionization-mode radio is set, ≥ 2 distinct non-Unassigned groups exist, and every assignable group has ≥ 2 samples (required for downstream statistics). The gate is **mode-agnostic** — Analysis Mode + KEGG species/Group are configured later on the Stage 3 setup screen, not here.

Missing values (NaN) vs true zeros (0.0)

metabolopan draws a hard line between a **measurement that is absent** and a **measurement that is genuinely zero**, and that distinction is carried — deliberately and consistently — through every downstream step. Imputing missing cells to 0 before analysis (a common habit) would silently bias the statistics, so this section spells out exactly what each value means and how it is treated.

The rule, fixed at load time. When an MS-DIAL .txt is parsed, an intensity cell that is empty / whitespace / "null" / "NA" (case-insensitive) or otherwise unparseable becomes `f64::NaN` — the internal marker for *missing / not measured / not computable*. A cell that literally reads `0` parses to a real `0.0`. Numeric metadata columns follow the same split: an empty cell is *absent* (`None`), while a written `0` is a real zero (and, because it would be a normalization divisor, is then rejected as a data-entry error rather than silently treated as absence).

The core behaviour: NaN is skipped, 0.0 participates. Every per-feature reduction in DAM — group mean, median, variance, IQR, and distinct-value count — first *drops* NaN values and computes on whatever remains. A `0.0`, being a real observation, enters the arithmetic in full. On the same three replicates this is the difference:

Group values	Effective <i>n</i>	Mean	Variance
[10, 12, NaN]	2	11.0	computed on 2 points
[10, 12, 0]	3	7.33	much larger

A missing replicate behaves as though that sample did not exist; a zero replicate pulls the mean down, inflates the spread, and counts toward the sample size.

Where the distinction surfaces, step by step:

- **Per-group pre-filter.** DAM requires ≥ 2 non-NaN values in *each* group. A NaN lowers that count — a group that is entirely NaN makes the feature un-testable, so it is skipped and tallied under *skipped* rather than occupying a "not significant" slot. A `0.0` counts as present and can help a group clear the minimum. (A group that is all identical zeros is instead removed by the "no variance" checks `nunique > 1` and `IQR > 0` — a different reason for the same skip.)
- **The statistical test.** Student / Welch / Brunner–Munzel each count non-NaN values and return a NaN *p*-value when a group has fewer than 2 after the NaN-drop. A `0.0` flows into the *t*-statistic, variance, standard error, and degrees of freedom like any other number.
- **Log transformation.** The optional variance-stabilising transform is `arcsinh (asinh)`, **not** `log10` — chosen precisely so zeros and small values are safe: `asinh(0) = 0` (a finite, usable value) whereas `log10(0) = -∞`. NaN is passed through untouched (the transform skips it; it never errors). This is a deliberate reason for preferring `arcsinh`: no pseudocount or clamping is needed to survive zeros.
- **Fold change.** Only a real `0.0` can drive a group mean (or median) to exactly 0 and so make `log2(FC) = ±∞`; those features are docked at the X-axis edges of the volcano plot (with a small jitter) and are never silently dropped. A NaN cannot produce this, because it was excluded from the mean in the first place — a NaN fold change is reserved for "the value genuinely could not be computed".
- **FDR correction.** Benjamini–Hochberg / Benjamini–Yekutieli skip NaN *p*-values entirely: they do **not** consume one of the *m* tests being corrected, and the NaN passes through to the output unchanged. A finite *p*-value — including one produced from real zeros in the data — is corrected normally.
- **Trend classification.** A feature whose adjusted *p*-value or `log2(FC)` is NaN is classified `NotSignificant`; it can never be called Up or Down.
- **NaN and ±∞ are kept distinct.** NaN means "could not be computed" (a group with $n < 2$; perfectly stratified groups under Brunner–Munzel). `±∞` is a real, *ordered* result — a *q*-value that underflows to exactly 0 becomes `+∞` on the `-log10` axis, and a zero-mean group gives a `±∞` fold change. Plots and CSV exports preserve the difference: the `+∞` point is docked at the plot edge, while the NaN point is dropped from the plot but still listed in the CSV.
- **CSV export.** A NaN cell is written as **empty** (`""`) — which round-trips back to "missing" if the file is re-read — while `±∞` are written as `inf` / `-inf` and a real zero is written as `0`.

One deliberate exception. PQN normalization treats a per-feature reference quotient of 0 the same as NaN: both are excluded as *unusable*, because a zero quotient carries no information for PQN's median-of-quotients factor. This is the only place the two are intentionally merged.

Bottom line. Leave a genuinely-missing measurement *empty* and write a measured zero as `0`; the software keeps them apart from input to export. If you pre-impute missing cells to `0`, you will inflate sample sizes, drag group means toward zero, distort variances and fold changes, and bias the differential-abundance calls — so let metabolopan carry missing values as NaN and do that bookkeeping for you.

Stage 2 — Sample normalization

Before any per-feature statistics, the user may select a *sample-axis* (column-wise) normalization to correct for technical variation between samples (injection volume, dilution, dry weight, total ion current). The matrix is normalized once at the start of every DAM run from the originally-parsed *intensity_raw*; *intensity_raw* is never mutated, so switching methods is lossless. The default is *None*, which preserves the prior behaviour bit-for-bit.

Five methods are offered in addition to the default:

- **Sum.** Per-sample factor = sum of non-NaN intensities for that sample. Output $x'[i, j] = x[i, j] / \text{sum}_j \times \text{median}_j(\text{sum}_j)$. Multiplying by the median of per-sample sums preserves overall magnitude so the Welch / Student path's optional *arcsinh* step (controlled by the Stage 2 Log transformation checkbox; default ON) stays in a useful range.
- **Median.** Same shape, using each sample's NaN-aware median as the factor.
- **Metadata column.** The user picks one of the optional numeric columns parsed from the metadata CSV (e.g. *dry_weight*, *dilution*). Each cell is divided by the per-sample value from that column, then rescaled by the median value across samples. Behavior on incomplete data:
 - ▶ *Missing value (empty cell):* the sample is **dropped** from the analysis — every cell in that sample's column is NaN-marked so DAM's NaN-aware machinery excludes it from per-feature statistics. The Stage 2 setup screen lists the samples being dropped in a yellow warning line before the user clicks "Start DAM".
 - ▶ *Non-positive value (zero or negative):* errors loudly with the offending sample and column named. Zero/negative metadata is a data-entry problem, not absence, so failing fast is the right call.
 - ▶ *Non-numeric cell:* parsed at CSV load time and errors before reaching Stage 2.
 - ▶ *Group preflight:* before any normalization work, the runner checks that dropping samples without a value still leaves at least 2 samples in each of the chosen numerator and denominator groups. If not, the error banner names the failing group, the column, the remaining count, and the minimum required (2).
- **Quantile.** Forces every sample's distribution onto a common reference (the per-rank mean across samples). Follows the **principle** Bolstad and Smyth converged on in the Bioconductor support thread #1569 (2003, <https://support.bioconductor.org/p/1569/>): tied items at sorted positions $[k, k+t]$ are assigned the **MEAN** of the reference values at those t rank positions — $\text{mean}(\text{reference}[k..k+t])$. This is the theoretical consensus they reached, and our code implements it literally. The widely-deployed canonical implementations (`preprocessCore::normalize.quantiles`, `limma::normalizeQuantiles`) both ship a simpler approximation — average-rank lookup with linear interpolation — which equals the principle only for $t == 2$ ties OR when the reference is locally linear, and diverges from the principle for $t \geq 3$ ties on a curvy reference (common at the bottom of below-LOD-imputed metabolomics samples, e.g. the worked example below). Our outputs therefore differ from `preprocessCore` / `limma` in exactly that case. Worked example: reference $[1.5, 7.5, 52.5, 502.5, 55000]$ with a 3-way tie at sorted positions 1–3 yields $\text{mean}(7.5, 52.5, 502.5) = 187.5$ here; `preprocessCore` / `limma` return $\text{ref}[2] = 52.5$. When all samples share the same non-NaN count, both schools produce identical numbers; the divergence is purely on unequal references.
- **Quantile — unequal non-NaN counts across samples.** When samples have different numbers of non-NaN cells (e.g. heterogeneous missingness), the reference is built on a common fractional-rank grid of size $K = \max(n_j)$ and each sample's sorted values are linearly interpolated onto it (matches `limma`'s $(r - 1) / (n - 1) \in [0, 1]$ scheme). This prevents the "longer samples dominate the high ranks" bug where a 3-non-NaN sample's largest value used to be mapped to the reference's 60th percentile (its `reference[2]` of 5 positions) instead of the reference's 100th percentile. When all samples share the same non-NaN count K , every fractional rank lands on an integer grid index, the interpolation paths collapse to direct lookups, and the output is bit-equal to the equal-length-only version we shipped before this change. NaN cells stay NaN.
- **PQN (Probabilistic Quotient Normalization).** Dieterle 2006: sum-normalize internally first; build a per-feature reference spectrum from the chosen cohort (default `ALL` samples, optionally restrict to one named group); for each sample compute the median of per-feature quotients vs the reference (skipping features where the reference is zero, NaN, or where the sample value is NaN); divide by that factor and rescale. Unassigned samples never reach this stage (they are dropped at the Stage 1 → Stage 2 boundary, so neither the reference cohort nor the per-sample factor loop sees them). If an *assigned* sample still produces a degenerate quotient median (NaN or 0), PQN aborts with `PqnDegenerateSamples` listing the offending names — switch to a different normalization method or remove the sample from the MS-DIAL .txt File type row. The dispatcher

INFO log line surfaces a `reference_features_used=N` field (added 2026-05-29) so you can see how many features the cohort actually anchored as PQN references (i.e. had `median(cohort) > 0`) vs the total feature count — useful for diagnosing QC sparsity without rerunning the pipeline.

Why Sum / Median / Metadata rescale to the median factor (rather than divide to a constant). These three methods share one driver. For each sample column j it computes a scalar factor f_j — the column sum (Sum), the column's NaN-aware median (Median), or the sample's positive metadata value (Metadata) — then rewrites every finite cell as

$$x'[i, j] = x[i, j] / f_j \times M, \quad \text{where } M = \text{median}_j(f_j)$$

The $\times M$ term is the deliberate part. The textbook form of sum normalization is plain x / f_j (or $\times 10^6$ for CPM-style counts), which forces every sample onto a *per-unit* scale: for Sum every column would then sum to 1 (proportions, $\sim 1e-5 \dots 1e-3$); for Median every column's median would become 1. We instead multiply back by M , the **median of the per-sample factors**, so each column's sum (resp. median) lands on M — the *typical* sample's original magnitude — instead of 1. The between-sample technical loading (injection volume, dilution, dry weight) is still equalized; only the absolute intensity scale is preserved.

- *Why it matters — the downstream arcsinh.* The default Log transformation is `arcsinh`, which behaves like a logarithm ($\text{arcsinh}(x) \approx \ln(2x)$) only once x is reasonably large; for x near 0 it is essentially **linear** ($\text{arcsinh}(x) \approx x$). Dividing to proportions would push the whole working matrix into that near-linear regime, collapsing `arcsinh`'s variance-stabilizing effect and degrading the t-test to a linear-scale comparison of tiny numbers. Keeping values at intensity scale ($\approx 1e4\text{--}1e7$) keeps `arcsinh` in its useful log-like regime — the "never near 0" goal. Scaling **all** cells by the same constant M cancels out in Brunner–Munzel's median ratio, but **not** under `arcsinh` (it is nonlinear), so this rescale specifically protects the Student / Welch + `arcsinh` path — the current default.
- *Why the median (not the mean) of the factors.* The median is robust — one unusually-high-loading sample cannot drag the target scale — and it makes the *typical* sample the anchor: that sample's $f_j \approx M$, so $f_j / M \approx 1$ leaves it essentially unchanged while the off-scale samples move toward it.
- *Worked numbers.* Three samples with column sums 6, 15, 24 give $M = \text{median}(6, 15, 24) = 15$; after $x / \text{sum}_j \times 15$ every column sums to **15** (sample A scaled $\times 2.5$, C $\times 0.625$, B unchanged) — not to 1. With Median normalization on per-sample medians 2, 20, 200, $M = 20$ and every column's median becomes **20**. Metadata is identical with f_j the chosen column's value, giving "intensity at the median dry weight."

Contrast with MetaboAnalyst, whose Sum/Median options divide to a fixed constant (or proportion) and pair it with `log10` (which additionally needs a pseudo-count to survive zeros); metabolopan's divide-then-rescale-to-median pairs with `arcsinh` so the normalization and the generalised-log transform stay numerically compatible. The chosen M is reported in the dispatcher INFO log as `scaling_to_median_factor=...`

Lifecycle. The normalization choice — and every other settings parameter — persists across every navigation transition for the lifetime of the session. Backing up to a previous stage never drops the choice; you simply land on the previous screen with all your prior picks intact. (If you re-pick files at Stage 1 and the prior numerator/denominator groups no longer exist in the new metadata, Stage 2 blocks the gate until you re-select valid groups.) There is no separate normalization step at Stage 3 — the working matrix (already normalized) is what Stage 3 enrichment sees.

Errors at startup. Normalization runs synchronously before the DAM tokio task spawns, so any failure (e.g. Sample 'A2' is missing a value in metadata column 'dry_weight') surfaces in the red banner immediately. The DAM task only starts when the working matrix is finite and shape-correct.

Caveats worth knowing.

- *Quantile* assumes the per-sample distributions *ought to* be the same. This is reasonable for replicate-rich studies of the same matrix (e.g. cell extracts) but breaks for cross-tissue or cross-organism comparisons where the biology genuinely differs at the distribution level.
- *PQN* is robust against most NMR-style dilution variation. The chosen reference cohort matters: when the study has a clean baseline group, using it as the PQN reference often produces sharper biological signal than **All samples**. **PQN is strict about sample quality**: a sample whose per-feature quotient median is NaN (no

usable features against the reference) or 0 (half-or-more of its non-reference-zero features are exactly 0 — typically a sparse / blank- like sample) is reported in an error message listing the offending sample names. Drop those samples from the metadata CSV or switch to a more tolerant method (None / Sum / Median / Metadata / Quantile). Pre-2026-05-26 these samples silently stayed un-normalized while their peers were PQN-scaled, producing biased differential-abundance calls from the scale mismatch.

- *Metadata* values MUST be strictly positive — division and the magnitude-preservation step assume positivity. Zeros and negatives error rather than silently passing through.
- *Sum/Median* preserve in-sample feature ratios exactly; they're scaled-to-magnitude versions of the same transformation. They differ in robustness: Sum is sensitive to a few high-intensity outliers per sample; Median ignores them.

Stage 2 — Deduplication by InChIKey

MS-DIAL routinely emits multiple Alignment IDs that resolve to the same compound. There are three biological / instrumental causes:

1. **Adduct multiplicity.** The same neutral molecule ionises as $[M+H]^+$, $[M+Na]^+$, $[M+NH_4]^+$, ... in positive mode (or $[M-H]^-$, $[M+Cl]^-$, $[M+FA-H]^-$, ... in negative mode). Each adduct yields its own Alignment ID but shares an InChIKey.
2. **Isotope peaks.** MS-DIAL emits separate rows for the M_0 monoisotopic peak and the $M+1$ / $M+2$ natural-abundance isotope peaks (flagged by Isotope tracking weight number or by an $[M+1]$ / $[M+2]$ suffix in Adduct type).
3. **Split chromatographic peaks.** When peak picking is suboptimal, a single Gaussian elution can be cut into two adjacent Alignment IDs that share every annotation but differ on Fill % / S/N average.

Feeding all duplicates into DAM inflates the FDR family size by 2–5× over the true compound count, eroding statistical power. Stage 3 ORA's κ (drawn-compound count) is also inflated, biasing pathway and module hypergeometric p-values toward whatever entries happen to contain commonly-multi-adducted compounds.

Deduplication runs as an opt-out toggle on the Stage 2 setup screen (default ON). The cascade is *purely* a deduplication operation, NOT a generic quality filter — features with `inchikey = None` pass through untouched, and singletons (one Alignment ID per InChIKey) are kept even if their annotation quality is poor.

Cascade decision table

Within each same-InChIKey group, the surviving feature is chosen by the first level of this cascade that distinguishes the two:

Level	Field	Rule
1a	MS/MS matched	True > False > blank
1b	Total score	larger wins (vendor-computed weighted composite of every spectral-similarity metric, incl. dot products)
2	Adduct class	Primary > NonPrimary > Dimer > Isotope; within Primary, $[M+H]^+$ / $[M-H]^-$ > $[M+Na]^+$ / $[M+NH_4]^+$ / $[M+K]^+$ / $[M+Cl]^-$
3a	Fill %	larger wins (per-sample peak coverage)
3b	S/N average	larger wins
4	Alignment ID	lexicographically smaller wins (deterministic terminator)

Adduct classification is deterministic and case-sensitive: Isotope is detected by either Isotope tracking weight number > 0 or an $[M+<n>]$ suffix in the adduct string; Dimer is detected by a leading multiplier ($[2M+H]^+$, $[3M-H]^-$, ...); Primary is the closed allowlist $\{[M+H]^+$, $[M+Na]^+$, $[M+NH_4]^+$, $[M+K]^+$, $[M-H]^-$, $[M+Cl]^-$ \}; everything else (including a missing adduct cell) is NonPrimary.

Audit CSV

The bottom-panel **Data** tab shows a "Download dedup audit (CSV)" button while on the Stage 2 result screen, whenever the DAM run was produced with dedup enabled (it is not shown on any post-Enrichment screen). The CSV format:

```
# Deduplication audit – generated by metabolopan
# Total dropped: <N>; total kept: <M>; null-InChIKey passthrough: <K>
dropped_alignment_id,inchkey,winner_alignment_id,decided_at,loser_value,winner_value
```

The `decided_at` column tells you which cascade level decided each drop (`MsmsMatched` / `TotalScore` / `AdductClass` / `FillPercent` / `SnAverage` / `Tiebreak`); `loser_value` and `winner_value` carry the deciding field's contents on each side (or empty when that side was `None`). In dual-mode runs the file contains one report per mode, separated by `# Mode: POS` / `# Mode: NEG` header lines.

Opt-out

Uncheck "Deduplicate features by InChIKey" on the Stage 2 setup screen to disable. With the box unchecked the DAM run is bit-equal to the pre-feature behaviour — every input row reaches the pre-filter, FDR m equals the post-pre-filter count, and `dedup_report` on the DAM result is `None`.

Stage 2 — DAM (Differentially Accumulated Metabolites)

Each feature is tested independently between the user-chosen numerator and denominator groups. Three statistical methods are offered; all follow the same overall flow.

1. Unknown-feature filter (default ON). Features whose `InChIKey` is `null` (MS-DIAL "Unknown" annotation; ~25% of the typical Alignment Result) are dropped before any statistical work, so the FDR correction's m does not include metabolites that cannot enter Stage 3 ORA anyway. The user can untick a "Drop unknown features (no InChIKey)" checkbox in Stage 2 Setup if they specifically want statistical results on unannotated features (e.g. to flag candidates for follow-up identification). This is a **deviation from the Python reference**, which keeps Unknown features in DAM and only drops them at the PubChem CID step — the user-visible toggle preserves the option to run Python-style.

2. Per-feature pre-filter. For every remaining feature, NaN values across the combined numerator u denominator columns are dropped, then we require, in order: (i) the numerator group has ≥ 2 non-NaN values, (ii) the denominator group has ≥ 2 non-NaN values, (iii) the combined `nunique` > 1 , and (iv) the combined `IQR` > 0 . Features failing any check are removed from the result and counted in a `skipped` tally visible in the UI. Checks (i) + (ii) were added 2026-05-29 — before this date, a feature where one entire group was NaN still passed the pre-filter and surfaced NaN through the test as an NS slot in the result; they are now skipped at the pre-filter layer instead so un-testable features no longer occupy NS slots downstream.

3a. Method: Student's t-test (equal variances) [parametric, **default**]. Classical (homoscedastic) form. Best when sample sizes per group are similar and the two groups have roughly comparable spread — under those assumptions it has slightly more power than Welch. **Default for new sessions:** paired with the `Log` transformation (`arcsinh`) step (also ON by default), it is the project's standard starting point. Switch to Welch if you suspect unequal variances, or to Brunner–Munzel when the distributions are skewed enough that the transform is not enough.

- Optional pre-test transform shared with Welch: when the Stage 2 setup `Log` transformation checkbox is checked (default ON; `SessionSettings.log_transform = true`), `arcsinh(x)` is applied to every non-NaN cell as the variance-stabilisation step (`asinh` handles zeros / negatives where `log10` would NaN them). When unchecked, this step is skipped and the raw working-matrix values flow straight into the t-test. The earlier hard-coded Pareto-scaling step was removed by `add-log-transform-and-scaling` (archived 2026-05-27) after empirical verification that per-feature linear rescaling cancels in the t -statistic — Welch / Student p -values under `log_transform=true` are bit-equal to the pre-change pipeline.
- Pooled variance $sp^2 = ((na - 1) \cdot va + (nb - 1) \cdot vb) / (na + nb - 2)$, fixed degrees of freedom $df = na + nb - 2$, two-tailed p via the Student- t CDF. Matches `scipy's ttest_ind(equal_var=True)`.
- **Fold change scale depends on log_transform.** Reason: under `log_transform=true`, the t statistic is computed on the `arcsinh`-transformed scale, but `arcsinh` is concave for positive values, so by Jensen's inequality the

raw mean ratio of two heavy-tailed groups can disagree in **sign** with the arcsinh-mean difference that the t -test actually evaluates. Reporting the raw mean ratio alongside an arcsinh-scale p value silently misclassifies outlier-driven features (e.g. $\text{num}=[0.1]\times 9 + [100]$ vs $\text{den}=[5]\times 10$ gives raw FC $\approx 2.02 \Rightarrow$ "Up", but Welch $t \approx -3.25$, $p \approx 0.01 \Rightarrow$ "Down"). Per the arcsinh-scale-fc-on-log-transform change (2026-05-29) the parametric arm now matches scales:

- ▶ `log_transform=false` (raw scale): $\text{FC} = \text{mean}(\text{numerator}) / \text{mean}(\text{denominator})$, $\log_2(\text{FC}) = \log_2(\text{FC})$. `FcBasis::Mean`. Bit-equal to the pre-2026-05-29 pipeline.
- ▶ `log_transform=true` (arcsinh scale): $\log_2(\text{FC}) = (\text{mean}(\text{arcsinh}(\text{num})) - \text{mean}(\text{arcsinh}(\text{den}))) / \ln(2)$, and $\text{FC} = 2^{\log_2(\text{FC})}$. `FcBasis::ArcsinhMean`. Sign of $\log_2(\text{FC})$ is **guaranteed** to agree with the t -statistic sign on the same data. For large x $\text{arcsinh}(x) \approx \ln(2x)$, so $\log_2(\text{FC})$ asymptotes to $\log_2(\text{GM}(\text{num}) / \text{GM}(\text{den}))$ — the classical log-fold-change of limma / DESeq2. For small x (close to 0) $\text{arcsinh}(x) \approx x$, so $\log_2(\text{FC})$ degrades to a scaled arithmetic mean difference rather than a true ratio. This is a documented consequence of variance-stabilisation; the equivalent log-FC interpretation only holds in the large- x asymptote where arcsinh aligns with $\ln$. CSV exports surface the active basis via the `fc_basis` column (mean / median / arcsinh-mean) so downstream consumers can identify which scale a number is on without rerunning the pipeline.

3b. Method: Welch's t-test (unequal variances) [alternative parametric]. Same parametric family as Student, but does not assume equal variances. Use this when group spreads visibly differ, or as a safer default when you are unsure.

- Same optional pre-test transform as Student (arcsinh only, gated by the Stage 2 Log transformation checkbox; default ON). Pareto scaling was removed for both parametric paths by `add-log-transform-and-scaling` (2026-05-27).
- Welch's t statistic is computed from the (optionally arcsinh-transformed) values using NaN-aware means and variances, with Welch–Satterthwaite degrees of freedom, then converted to a two-tailed p via the Student- t CDF. Matches `scipy's ttest_ind(equal_var=False)`.
- **Fold change scale matches the test scale** — same rule as Student above. With `log_transform=true`, FC is on the arcsinh scale (`FcBasis::ArcsinhMean`), so its sign always agrees with the Welch t sign. With `log_transform=false`, FC is the classical raw mean ratio (`FcBasis::Mean`).

Welch / Student edge case — zero variance in one group. When every replicate in one group has the same value (e.g. the feature is below the limit of detection in every sample of one condition and gets imputed to a constant), the Welch–Satterthwaite degrees of freedom collapse to $n - 1$ of the OTHER group. For $n = 2$ this gives $\text{df} = 1$, which makes the t distribution extremely wide and the p -value very conservative — even when the two groups are visibly well-separated. This is the standard mathematical behaviour (matches R's `t.test(var.equal=FALSE)` and SciPy's `ttest_ind(equal_var=False)` exactly), but for metabolomics the affected features often correspond to genuine "presence-in-one-condition, absence-in-the-other" signals you may want to keep. `run_dam` emits a single INFO log per run reporting the count of features that triggered this path (look for `zero_variance_features=N` in your session log at `<binary>/data/logs/session_*.log`); when $N > 0$, consider re-running with the Brunner–Munzel method, which is rank-based and handles this edge case differently. As of 2026-05-29 the diagnostic counter uses a relative tolerance — variance below $(\max(|\text{mean}|, 1))^2 \times 1\text{e-}20$ is flagged — so features whose group is constant up to floating-point noise (e.g. arithmetic on bit-equal pre-norm inputs at high intensity scale where $\text{var} \approx \epsilon^2 \times c^2$ is non-zero but the df pathology kicks in the same way) also contribute to the count. The per-method `var == 0.0` guard inside the t -test functions themselves is unchanged — only this diagnostic counter was loosened.

3c. Method: Brunner–Munzel + Cliff's δ [non-parametric]. Appropriate when the intensity distributions are skewed or unequal across groups and the variance-stabilising transform is not enough. Metabolomics data is often poorly described by Gaussian assumptions (highly skewed log-distributions, frequent presence/absence patterns, batch artefacts), so in those cases Brunner–Munzel + Cliff's δ can deliver more honest p -values across the kind of spreads the workflow sees. Select it via the Stage 2 setup radio when the default Student's t -test (even after arcsinh) is a poor fit — e.g. heavily skewed or presence/absence-dominated features, or when matching a previously published non-parametric analysis.

- The Brunner–Munzel statistic is computed with mid-ranks across numerator u denominator, combined with Welch–Satterthwaite-like degrees of freedom, then converted to a two-tailed p value via the Student- t

distribution. Behaviour matches SciPy's `brunnermunzel(distribution='t')` and R's `lawstat::brunner.munzel.test` — the W denominator inside `sqrt` is $n_x \cdot S_x + n_y \cdot S_y$ (the pre-2026-05-26 implementation used $(n_x + n_y) \cdot (S_x / n_x + S_y / n_y)$, which inflated $|W|$ by $\sqrt{N/2}$ at equal n ; pre-fix BM p-values were systematically $\sim 1.58\times$ too significant at $n=5$ per group. Prior CSV / volcano exports under BM should be regenerated.).

- Cliff's δ effect size: $(gt - lt) / (n \cdot m)$ where gt and lt are the strict-greater and strict-less pair counts. Range $-1 \dots +1$; $|\delta| \geq 0.33$ is the conventional "medium effect" threshold used here.
- Fold change uses group **medians**: $FC = \text{median}(\text{numerator}) / \text{median}(\text{denominator})$ and $\log_2(FC) = \log_2(FC)$. Medians are robust against outliers, matching the rank-based-test philosophy.

4. Multiple-testing correction. Each Stage 2 run applies a user-selected FDR correction to the per-feature p values, regardless of which statistical method produced them. The Stage 2 setup screen exposes a radio with two options:

- **Benjamini–Hochberg (BH)** — the default. Matches R `p.adjust(method='BH')` and MetaboAnalyst conventions, which makes Stage 2 results directly comparable to those tools' published numbers. BH assumes independence or positive regression dependence between tests.
- **Benjamini–Yekutieli (BY)** — opt-in. Multiplies BH's q-values by the exact harmonic factor $c(m) = \sum_{i=1}^m 1/i$ ($\approx \ln(m) + \gamma$ for large m , so BY is roughly $9\times$ more conservative than BH at $m = 5,000$). BY controls FDR under arbitrary positive dependence, so it is the safer choice when many features are biologically correlated (e.g. metabolites sharing pathway membership).

NaN p values pass through corrected as NaN under either method. The chosen method is reported on the volcano annotation strip (e.g. `FDR(BH)<0.05`) and written as a leading `# FDR: BH / # FDR: BY` comment line on every DAM CSV export, so screenshots and downloads are self-documenting. References: Benjamini & Hochberg (1995); Benjamini & Yekutieli (2001).

5. Trend classification (recomputed live as the user adjusts thresholds — never stored in the result). Default thresholds: $FC = 2.0$ (equivalently $|\log_2(FC)| \geq 1.0$), $FDR = 0.05$, $|\delta| \geq 0.33$ (BM only).

- Student / Welch (both parametric, no effect size): Up iff $FDR < \text{threshold AND } \log_2(FC) > \log_2(\text{fc_threshold})$; Down iff $FDR < \text{threshold AND } \log_2(FC) < -\log_2(\text{fc_threshold})$. The δ threshold is ignored for the parametric tests.
- Brunner–Munzel: the parametric rule **AND** $|\delta| \geq \text{delta_threshold}$. Features with $\delta = \text{None}$ (BM was unable to compute the effect size because one group had fewer than 2 non-NaN values) are classified `NotSignificant`.

6. Volcano plot. X axis = $\log_2(FC)$, Y axis = $-\log_{10}(p_adjusted)$. **What the X axis represents depends on the active method and the Log transformation toggle** — mean ratio for Welch / Student with `log_transform=false`, arcsinh-mean difference (in $\log_2$ units) for Welch / Student with `log_transform=true`, median ratio for Brunner–Munzel. See section 3 above; the active basis is recorded on each `DamFeature` as `fc_basis` (mean / arcsinh-mean / median). Three colours follow the trend classification (red / blue / grey, $\alpha \approx 0.5$). Threshold lines are dashed black: the horizontal line at $-\log_{10}(FDR)$, the vertical lines at $\pm \log_2(FC)$. Features whose $\log_2(FC)$ is $\pm\infty$ (one group's mean or median is exactly 0) are docked at the X-axis edges $\pm(xabs_max + 0.5)$ with a small jitter so they stay visible. Symmetric saturation on the Y axis: features whose BH/BY q-value underflows to exactly 0.0 (very large $|t|$ / very small raw p, common for well-separated groups) are docked **just below** the Y-axis top (`y_max`) with a per-point downward jitter of up to 0.08 in $-\log_{10}(q)$ units (scale-matched to the X-axis ± 0.04 jitter convention) so multiple saturated features don't pile at a single pixel (added 2026-05-29 — before this date all $q=0$ features mapped to exactly `y_max` and visually collapsed). The underlying `neg_log10_p_adjusted` value is still `f64::INFINITY`, recorded as such in the CSV export — only the on-screen position is jittered. The Y axis is otherwise clipped at `finite_max + 1` for display only; underlying numeric values are preserved in the CSV export. NaN `neg_log10_p_adjusted` is reserved for genuine "p couldn't be computed" cases (BM perfectly stratified groups; parametric test with $n < 2$ after NaN-drop) — those points are dropped from the plot but still listed in the CSV. A single annotation strip below the X-axis label summarises method, the active FC basis (FC: mean / FC: median / FC: arcsinh-mean, added 2026-05-29 so the X-axis interpretation is unambiguous without consulting the CSV), active thresholds, and $\pm\infty$ counts — e.g. `Method: Brunner-Munzel | FC: median | FDR(BH)<0.05, FC≥2.0, |δ|≥0.33 | -∞: 12 +∞: 8`.

BM dot size encodes Cliff's δ magnitude. On Brunner–Munzel renders, each scatter dot's radius is mapped from the feature's `|Cliff's δ|`: $|\delta|=0$ gives the smallest still-visible dot, $|\delta|=1$ gives a dot $\approx 1.3\times$ the default radius, and intermediate magnitudes scale linearly between the two anchors. The right-side legend grows a second |

δ size block under the existing trend counts, with three reference dots at $|\delta|=0/0.5/1.0$ in neutral grey — size-match scatter dots against these references to read magnitude off the chart. Welch / Student renders keep a uniform dot radius across the chart and do NOT draw the $|\delta|$ size legend section (those tests don't produce a Cliff's δ to encode). BM features where $|\delta|$ is undefined (one group with $n < 2$ non-NaN values) fall back to the default radius and still render in the appropriate trend colour.

Stage 2 caveats worth knowing.

- BM's median-based FC means small-n studies (e.g. 3 samples per group) are more likely to produce $\pm\infty \log_2(\text{FC})$ than Welch's mean-based FC, because a single zero in three samples drives the group median to zero. The annotation strip surfaces the $\pm\infty$ count so this is never silent.
- A parametric t-test (Student or Welch) on $n = 2$ per group has only $\sim 1-2$ degrees of freedom and is unreliable; the Stage 1 gate's ≥ 2 samples-per-group requirement keeps you above the floor, but well-powered parametric tests want ≥ 5 per group. Student with equal sample sizes is the most sensitive of the three when its equal-variance assumption holds; Welch is the robust fallback when variances visibly differ.
- Trend classification depends on the active thresholds. The CSV exporter writes the trend computed at export time, which is identical to what the volcano shows at that moment.

Stage 3 — Enrichment (over-representation analysis)

Stage 3 takes the DAM result from Stage 2 and asks: "*Which KEGG entries are over-represented in my list of differentially-abundant compounds?*" — where "entry" means a **KEGG pathway** (pathway mode) or a **KEGG module** (module mode). The two modes share identical machinery for the hypergeometric test, user-selected FDR (BH or BY), and the measurable-metabolome universe; they differ only in what catalogue of compound-sets ORA operates over.

Stage 3 carries its own FDR-correction radio, independent of the Stage 2 choice — the default is BH for cross-tool reproducibility, but **BY is the recommended choice for pathway/module ORA**, because entries inherently share compounds (many cpds appear in multiple pathways), which violates BH's independence assumption. The Stage 3 dot plot's colorbar title and annotation strip both name the active method (e.g. $-\log_{10}(\text{FDR (BY)})$) / FDR: BY), and the leading # FDR: comment line on the enrichment CSV records the choice for downstream parsing.

Enrichment Analysis setup screen

The Stage 3 setup screen is where the user picks:

- **Analysis Mode** (Pathway / Module) via a radio toggle. Both modes' selections AND their fetched KEGG caches coexist for the lifetime of the session — toggling between modes is instant and never re-fetches data you've already pulled.
- **KEGG scope**. Pathway mode shows a searchable species selector with the eagerly-loaded KEGG organism list; module mode shows the Level + Group picker described below in *Module mode*. Selecting a species (or Group) triggers the corresponding KEGG fetch inline on this screen — a small progress bar with caption streams the per-pathway (or per-module + ETA) progress without leaving the setup screen.
- **Direction** to include (Both / Up only / Down only).
- **Minimum number of compounds detected in a pathway/module** (the "minimum entry size" filter; the label is mode-aware — "... in a pathway" in Pathway mode, "... in a module" in Module mode; default 1, range [1, 20]). Drops pathways / modules whose **universe-restricted** compound count ($m_p = |\text{entry.comounds} \cap \text{universe}|$) is below this threshold **before** the FDR family is built. Both m_p (here) and the hypergeometric m parameter use the **set** cardinality of the intersection: a compound listed more than once in a KEGG entry's COMPOUND block is counted **once**, not per-occurrence — counting raw occurrences would inflate m (and the expected hit count) and deflate the entry's enrichment ratio (a duplicate-counting bug fixed 2026-05-26). The default 1 is permissive — only $m_p = 0$ entries are dropped (they short-circuit to $p = 1.0$ anyway), so every pathway with at least one measurable compound is tested. Raise it to **3 to match MetaboAnalyst's minPathSize** and exclude the small-but- mathematically-untestable $m_p \in \{1, 2\}$ entries — an entry with $m_p = 1$ can never produce a BH-significant result no matter how strong the hit, so dropping them reduces the multiple-testing penalty (at the cost of testing fewer pathways). Orthogonal to *Minimum hit count*: this knob

is a **pre-FDR ENTRY filter** that shrinks m ; *Minimum hit count* is a **post-FDR DISPLAY filter** that doesn't change p-values.

- **Enrichment FDR threshold** (default 0.05).
- **FDR correction** (BH default; BY recommended for ORA — see above).
- **Minimum hit count** (post-FDR display filter; default 1).
- The **Run Enrichment** button (disabled while a fetch is in flight; the disabled-state hover tooltip explains which fetch is blocking the button). The Top N input that controls the dot plot's display cap lives on the next screen (Enrichment Result) so you can iterate on it after seeing the data, without coming back to setup.

Pathway mode

The pipeline is:

1. **Identity resolution (PubChem PUG REST).** For every feature passing Stage 2's pre-filter (NOT just DAM-significant ones), resolve its InChIKey to one or more PubChem CIDs via a POST to `compound/inchikey/property/InChIKey/CSV`. Up to 200 InChIKeys per batch.
2. **KEGG compound conversion (KEGG REST).** For every unique CID, resolve to a KEGG compound (`cpd:Cxxxxx`) via `/conv/compound/pubchem:CID1+CID2+...`. Up to 10 CIDs per batch, throttled by the KEGG client (334 ms between requests, ~3 req/s under KEGG's documented soft cap). Lines that map to `glycan:` or `dr:` are filtered out — only `cpd:` targets are kept. HTTP 403 is treated as a rate-limit signal and retried up to 5× with 5 s backoff.
3. **Multi-mapping rule.** One feature is one chemical entity. If PubChem returns multiple CIDs for an InChIKey (stereo / regio / salt ambiguity) and they all resolve to the same KEGG `cpd`, the feature contributes that `cpd` **once** to the DAM compound set K and to the universe N . If they resolve to genuinely different `cpds` (very rare for sane PubChem records), each `cpd` counts. Features whose InChIKey has no PubChem CID, or whose CIDs all fail to map to a KEGG `cpd`, are dropped from K and N and surfaced in the bottom-panel **Data** tab's mapping funnel ($\langle N \rangle$ InChIKeys $\rightarrow$ $\langle N \rangle$ PubChem CIDs $\rightarrow$ $\langle N \rangle$ KEGG `cpds`).
4. **Universe definition (N).** The universe is the union of unique `cpd` IDs across all annotated features that passed Stage 2's pre-filter AND successfully mapped through PubChem and KEGG conv — the *measurable metabolome* on this platform. This is more conservative than MetaboAnalyst's "all KEGG compounds in this species" universe, which inflates N with compounds your instrument cannot detect. We intentionally use the measurable-only universe so p-values better reflect what your data could have said.
5. **Pre-FDR entry-size filter.** Before any hypergeometric work, each pathway's M_p is compared against the user-tunable `min_entry_size` (default 1, range [1, 20]). Entries with $M_p < \text{min_entry_size}$ are **dropped entirely** from the run — they contribute no p-value to the FDR family, do not appear in the CSV, and do not appear on the dot plot. The dropped count is surfaced in the bottom-panel **Data** tab via a retention line Tested: $\langle \text{surviving} \rangle / \langle \text{total} \rangle$ ($\geq N$ compounds in universe) (Tested: $\langle \text{surviving} \rangle$ ($\geq N$ compounds in universe) in module mode). The default 1 keeps the pre-filter permissive — only $M_p = 0$ entries are dropped (they were short-circuiting to $p = 1.0$ anyway), so every pathway with at least one measurable compound is tested. Raising it to **3 matches MetaboAnalyst's minPathSize** and additionally excludes entries with $M_p \in \{1, 2\}$ that are mathematically untestable at typical K/N values — e.g. an entry with $M_p = 1$ can produce at most $k_p = 1$, giving raw $p \approx K/N$ which is rarely below $\alpha = 0.05$ and even more rarely below the BH critical value $0.05/m$. The trade-off is symmetric: a lower `min_entry_size` tests more pathways but enlarges the multiple-testing family m .
6. **Per-pathway hypergeometric test.** For each pathway p surviving the entry-size filter, with $M_p = |\text{unique}(\text{pathway.compounds}) \cap \text{universe}|$ (set cardinality of the pathway's unique `cpd` IDs that fall within the measurable universe — duplicate `cpd` IDs within a single COMPOUND block do NOT inflate M_p) and $k_p = |\text{pathway.compounds} \cap \text{universe}|$:
 - $p_value = 1 - \text{HypergeometricCDF}(k_p - 1; N, M_p, K)$ (upper-tail probability of seeing AT LEAST k_p hits)
 - If any of k_p , M_p , K , N is zero, the implementation short-circuits to $p_value = 1.0$ (avoids undefined CDF arguments).
 - The implementation also enforces $K \leq N$ as a `debug_assert!` (any upstream regression that lets K leak compounds outside N is caught loudly in dev/test; release builds emit a per-run ERROR log summarising

any Hypergeometric domain errors so the failure mode "all entries non-significant with no diagnostic" cannot ship silently).

7. **User-selected FDR correction** via the Stage 3 setup screen's independent radio (default Benjamini-Hochberg; Benjamini-Yekutieli one click away; **None** as a third option for exploratory runs only, see below). The radio is independent of the Stage 2 choice on purpose (per the archived `add-bh-fdr design.md D3`): the two stages have different dependence profiles, and many users will reasonably want Stage 2 BH (cross-tool reproducibility on the volcano) + Stage 3 BY (conservative ORA on shared-compound entries). For pathway/module ORA we **recommend BY** even though MetaboAnalyst and most biology tools default to BH: pathways share compounds heavily (glycolysis ↔ TCA share G6P, pyruvate, etc.), so the independence assumption underlying BH is violated. BY is more conservative under dependence; expect uniformly higher (less significant) adjusted p-values than MetaboAnalyst on the same data. **None** skips multiple-testing correction entirely — the `fdr` field in the result table and CSV carries the raw p-value verbatim. Use **only for exploratory ranking**, never for published claims of significance; on a typical KEGG pathway catalogue (~300 pathways tested) you'd expect ~15 false positives at $p < 0.05$ purely by chance. The Stage 2 DAM radio does NOT expose None — raw p across ~13 k features would flood the result set; a hand-crafted snapshot carrying `dam_fdr_method=NoCorrection` is defensively coerced back to BH with a `tracing::warn!` event. **Colour scale.** Each marker's fill encodes $-\log_{10}(\text{FDR})$ (raw $-\log_{10}(p)$ under None) on a ColorBrewer **YlOrRd** 9-step ramp — palest yellow for the least-significant entry shown (FDR at the displayed threshold) deepening to dark red for the most significant; the dots and the colorbar legend share a single $-\log_{10}$ span, so equal colours mean equal significance across both. The active method is recorded in the dot plot's colorbar title ($-\log_{10}(\text{FDR (BH)}) / -\log_{10}(\text{FDR (BY)}) / -\log_{10}(p\text{-value})$ for None — the wrapper drops because the axis values ARE raw p, not q) and in the leading `# FDR: BH / # FDR: BY / # FDR: None` line of the exported enrichment CSV. The CSV also carries a second comment line `# MinEntrySize: N` recording the pre-FDR filter threshold used in that run, so the file is self-documenting. The dot plot itself also carries a four-line plain-language annotation block below the X axis so reviewers can reconstruct the FDR family from the figure alone: Background universe = <N> compounds measured and mapped to KEGG / Compounds of interest = <K> differentially abundant (increased | decreased | both directions) / Pathways tested = <m>[of <total> · <dropped> skipped (< <min_entry> compounds each)][; ≥ <min_hit> hits required] / Significance: FDR-adjusted, Benjamini-Hochberg (BH) (or ... Benjamini-Yekutieli (BY), or raw p-value (no FDR correction)). The N / K / m symbols are deliberately spelled out rather than abbreviated; the tested count <m> is the number of entries that reached BH/BY and is the divisor each raw p-value was multiplied by. The m denominator equals the count of pathways that **survived the pre-FDR min_entry_size filter** (step 5) — i.e. `m = entries.len() - entries_dropped_by_min_entry_size`. Pre-filter entries with `M_p = 0` are also dropped under the default `min_entry_size = 1`. The orchestrator-level Group filter (module mode) is applied at an even earlier layer; m reflects both filters by the time FDR runs.
8. **Display filtering (post-FDR).** A user-controlled `min_hit_count` (default 1) hides pathways with fewer hits from the dot plot and CSV. This is a *display* filter — m was already computed over all surviving entries, so the FDR values are honest regardless of this setting. Distinct from `min_entry_size` in step 5: that one is a **pre-FDR ENTRY filter** that shrinks m; this one is a **post-FDR DISPLAY filter** that doesn't change p-values.
9. **Dot plot selection vs ordering — two different bases.** The dot plot chooses *which* entries to draw and *how* to stack them on the Y axis using **deliberately different criteria**:
 - **Selection (which entries appear) is by statistical significance.** Among entries passing `fdr < threshold` and the `min_hit_count` filter (steps 7–8), the plot keeps the **Top N with the lowest FDR** (`top_n`, default 20, tunable on the result screen). The entries shown are therefore always the *most significant* ones — they are **never** selected by fold enrichment.
 - **Vertical order (Y axis) is by effect size.** The kept entries are then arranged by **fold enrichment (observed/expected) descending**, so the entry with the largest fold enrichment sits at the **top** and the figure reads as a largest-on-top staircase down the X axis (which is itself fold enrichment). Ties break by FDR (more significant first), then entry ID. This matches the clusterProfiler convention of ordering the Y axis by the X-axis metric.

The practical consequence: when more entries are significant than `top_n`, the ones omitted are the **least significant** (highest FDR) — *not* the smallest fold enrichment. Significance gates inclusion; effect size only arranges what got in. The exported CSV is independent of this: it lists every surviving entry ordered by ascending FDR, with full (untruncated) names.

10. **Dot plot canvas height.** The exported plot height auto-fits to the number of rows actually shown — $\text{clamp}(\min(\text{top_n}, \text{displayed}) \times 0.3 + 1.0, 2.0, 40)$ inches — and is **recomputed every time you Draw / Re-draw**. So if a run is non-significant at your initial FDR threshold and you loosen the threshold on the result screen and redraw, the canvas grows to fit the newly-revealed rows instead of cramming them into a short plot (which would truncate the Y-axis labels). Editing the **Height (in)** field turns it into a manual override that sticks until the next enrichment run/re-run resets the auto-fit.

Text size is independent of entry count. Labels, axis titles, the colorbar, and the Hits legend scale with the plot **width** (a fixed `Width (in) × DPI`), *not* the auto-fitting height — so a two-entry result renders its text at exactly the same size as a twenty-entry one. (Before this was fixed, sparse plots keyed their font size off the short canvas and came out with tiny, hard-to-read text.) The 2.0-inch lower bound on the height — raised from 1.5 — exists so the full-size legend always clears the canvas on those sparse results.

Module mode

Module mode runs the identical PubChem → KEGG conv → hypergeometric → user-selected-FDR pipeline as pathway mode, but **(a)** the entry catalogue is the set of KEGG modules instead of per-species pathways, and **(b)** the user picks an **organism Group** instead of a single species. A module is included in the analysis when its KEGG COMPLETE block contains at least `min_group_overlap` organisms from the chosen Group; this is how the per-species framing maps onto the global module catalogue.

1. **Organism Group selection.** The Stage 3 **Enrichment Analysis setup** screen surfaces a Level radio (1 / 2 / 3) and a Group dropdown when the Analysis Mode toggle is set to Module. The Level indexes into the KEGG / list/organism lineage column (Eukaryotes at Level 1, Animals / Bacteria / etc. at Level 2, Mammals / Insects / etc. at Level 3). KEGG currently has 11,744 organisms, all with exactly 4 lineage levels; we expose the first three. Picking a Group materialises `org_codes`: the set of KEGG organism codes (`hsa`, `ath`, ...) belonging to that Group. (Prior to the `reorder-gui-and-move-mode-to-stage3` change archived 2026-05-25 this selector lived on Stage 1; it now lives on Stage 3 setup so you only commit to a mode after seeing the DAM result.)
2. **Module → Group filter.** Each module's `/get/<module-id>` response carries a COMPLETE block listing the organisms in which the module is fully assembled. A module is retained for ORA when:

$$|\text{module.complete_orgs} \cap \text{group_orgs}| \geq \text{min_group_overlap}$$

The default `min_group_overlap = 1` is permissive ($\exists$ -overlap: any single organism in the Group is enough). Higher values tighten the filter — e.g. `min_group_overlap = 5` requires that at least 5 of the Group's organisms have the module fully assembled. The active threshold is surfaced in the Stage 3 result header so any number you publish is reproducible from header + cache snapshot alone.

3. **Universe and K — same as pathway mode.** The PubChem and KEGG-conv phases are mode-agnostic. `N` is still the measurable metabolome (DAM features that mapped through to a KEGG cpd); `K` is still the cpd set of DAM features matching the active direction filter (Up / Down / Both). Module mode does *not* substitute "all module compounds" or "all KEGG compounds" for `N`.
4. **Per-module hypergeometric test.** Identical to pathway mode: for each retained module `m`, $M_m = |\text{module.compounds} \cap \text{universe}|$, $k_m = |\text{K} \cap \text{module.compounds}|$, and $p_value = 1 - \text{HypergeometricCDF}(k_m - 1; N, M_m, K)$ with the same zero-input short-circuit.
5. **User-selected FDR correction** — same options and defaults as pathway mode (BH default; BY recommended for shared-compound entries). The `m` denominator equals the count of **retained modules** (after the Group filter), not the total ~573 modules in the KEGG catalogue. This is the correct null: ORA is asking "among the modules that *could* apply to this organism Group, which are over-represented?" Including taxonomically-irrelevant modules in `m` would distort the FDR upward without contributing biological signal.

6. **Empty-COMPOUND module counter.** Some KEGG modules (signature / reaction-only modules) have no COMPOUND block at all. Their entries pass through ORA with `compounds = []` and short-circuit to `p_value = 1.0`. The bottom-panel **Data** tab surfaces these as a `With compound list: <kept> (-<empty> empty)` line so silent drops never erode trust. (Symmetrical pathway-mode reporting is on the roadmap.)

Module-mode caveats worth knowing.

- **First-run cost.** A cold fetch of all ~573 currently-listed modules from KEGG takes approximately 6–8 minutes at the 334 ms inter-request throttle (3 req/s). The module ID range is M00001–M01063 but KEGG keeps the range sparse — retired IDs aren't reused, so the actual count is lower than the upper bound. The Stage 3 setup screen shows an inline progress bar with an ETA derived from a rolling-average of per-module wall-clock time once the first 5 modules have completed. Subsequent runs use the cache and the Run Enrichment button enables in seconds.
- **Group 1 has only two options** (Prokaryotes / Eukaryotes), which is biologically very coarse. It exists for completeness — e.g. "any prokaryote" comparative studies — but most analyses will benefit from Level 2 (6 candidates) or Level 3 (tens of candidates) for finer scoping.
- **min_group_overlap is a research knob.** Default 1 (permissive $\exists$ -overlap) is appropriate for exploratory work. For papers, consider testing a higher threshold to ensure robustness — a module that only one of the hundreds of organisms in a Group (e.g. "Animals") possesses is biologically marginal for that analytic frame even if it survives the default filter.
- **Module CSV column names match pathway-mode CSV.** Both modes export the same header: `EntryID,EntryName,Hits,Total,Expected,EnrichmentRatio,PValue,FDR,HitKeggIDs`. In module mode the `EntryID` column carries M00001-style module IDs; in pathway mode it carries `<species_code><pathway_number>` IDs (e.g. `gmx00010`).

Starting a new analysis round

When you finish an enrichment run and want to analyse a different dataset — or re-run the whole pipeline from scratch — the Stage 3 **Enrichment Result** screen offers a **Start a new analysis** button on its own line below [Download enrichment results CSV]. Clicking it opens a confirmation dialog warning that the current DAM / enrichment results and any un-downloaded plots or CSV will be lost. On **Start over** the app resets every parameter to its default, clears the loaded MS-DIAL .txt / metadata .csv and the in-memory KEGG data, and returns you to Stage 1 — *without* re-running the startup organism-list load. (The on-disk KEGG cache survives, so re-fetching the same species or modules afterward is a fast cache hit.)

This is deliberately distinct from the stage stepper's **Input** step, which navigates back to Stage 1 while *preserving* every setting, loaded file, and fetched cache so you can keep iterating on the **same** dataset. Use the stepper to tweak and re-run the current analysis; use **Start a new analysis** to discard everything and begin fresh. If you might want the current configuration again, save it via the Data tab's [Save settings...] button before starting over.

Dual-mode (positive + negative ionization) input

Metabolomics experiments often run the same biological samples through both positive and negative ionization modes, producing two MS-DIAL .txt exports per study. The app supports loading both files at once and combining their enrichment signal under a conservative union rule.

When to use dual-mode

Use dual-mode whenever you have BOTH a POS and NEG .txt for the same biological samples and want a single enrichment result that reflects evidence from either ionization. Single-mode (one .txt) remains the default. (Since 2026-06-01, single-mode K also applies the conflict-only-strict rule — see *Stage 3 — N and K math* below — so it is bit-equal to prior releases EXCEPT for datasets where a compound is reached by both an Up and a Down feature, which is now excluded.)

Preparing inputs

1. **Two .txt files.** One per ion mode. The `Adduct` type column drives BOTH the auto-fill of slot 1's mode radio (see *Stage 1 UI* below) AND an advisory disagreement hint when the user manually overrides to the opposite polarity (adducts ending in + infer Positive, in - Negative).

2. **One metadata CSV that includes a biosample column** (e.g. header `sample,biosample,group`; column order is free). Each row maps a per-mode sample name (e.g. `CTR_positive_01`, `CTR_negative_01`) to its **biosample label** (the same `CTR-01` for both modes) and group. The biosample column lets the tool recognize that two differently-named samples are the same biological replicate.

A dual-mode run with a CSV that has no biosample column is blocked at Stage 1 with a specific error — add the biosample column or remove the second .txt to proceed.

Single-mode does NOT need a biosample column. It is only required when a second .txt is loaded. With one .txt, the plain `sample,group` form is enough (a biosample column, if present, is recognized by name and excluded from the Stage 2 metadata-normalization radio — it is not offered as a numeric metadata column).

Unbalanced or missing-mode samples

The worked example below is perfectly balanced (every biosample runs in both modes), but real studies sometimes acquire a biosample in only one polarity. The biosample column is what pairs the two modes, so Stage 1 enforces three dual-mode integrity checks **before** Continue to DAM is allowed. Each surfaces a specific error:

1. **Each group needs ≥ 2 samples in each mode.** If a group has enough replicates in POS but drops below 2 in NEG (e.g. because several biosamples lack a NEG acquisition), Stage 1 blocks with Group 'X' has N sample(s) in POS but M in NEG — both modes need ≥ 2 . A few missing-mode samples are tolerated as long as every group still clears 2-per-mode; it is only the per-group, per-mode count that gates.
2. **A biosample must be unique within a mode.** Two rows mapping the same biosample label to the same mode trip Biosample 'B' appears in N POS rows — must be unique per mode.
3. **A biosample must stay in the same group across modes.** If `CTR-01` is control in POS but treatment in NEG, Stage 1 blocks with Biosample 'B' is in group 'X' in POS but 'Y' in NEG.

Effect of a missing-mode sample that passes the gate. The two modes run DAM independently on their own sample columns — a biosample absent from NEG simply isn't iterated in the NEG run, so that mode has fewer replicates for its group and correspondingly lower statistical power; it does not invalidate the run. At Stage 3 the union is built at the **compound** level (per the conflict-only-strict rule below), not the sample level, so a biosample missing one mode just lets that mode contribute Absent for the affected compounds — the integrated K is unaffected.

Recommendation. For the cleanest dual-mode result, acquire every biosample in both polarities. If some samples are genuinely single-polarity, either keep them only in the mode where they exist (as long as each group still has ≥ 2 per mode), or drop the unbalanced side. Samples that appear in a .txt but not in the metadata CSV are flagged Unassigned and auto-dropped at the Stage 1 → Stage 2 boundary (see the group-mapping notes above), which is another way to exclude an unwanted column.

Stage 1 UI

Slot #1 (always visible) and slot #2 (revealed by the + Add second ionization mode button) each have a file picker, a mode radio (Positive / Negative), and a per-slot summary. After `auto-infer-stage1-ion-mode` (archived 2026-05-26) the slot-1 mode radio auto-fills from `infer_polarity(&table)` on every fresh file load and re-pick: a $\geq 95\%$ positive-suffix Adduct column sets it to Positive, $\geq 95\%$ negative-suffix sets it to Negative, ambiguous mixtures leave the radio unset (the existing grey "Could not auto-detect..." hint still applies). When slot 1's mode is set, slot 2's radio auto-fills to the **opposite** on three triggers: (1) slot 2 is revealed via the + Add second ionization mode button, (2) slot 2's .txt is loaded, (3) slot 1's mode changes (manual click or re-pick re-inference) — case (3) also flips slot 2 if the new slot-1 value collides with what slot 2 already showed. The user may still manually click any radio to override. The slot-2 radio still disables the option already chosen for slot #1 (a tooltip explains why). The adduct-disagreement hint ("yellow: Adduct column says X but you selected Y") still fires on a manual override that contradicts auto-detection; neither hint blocks Continue to DAM. The Stage 1 screen does NOT carry an Analysis Mode toggle or any KEGG selector — those live on the Stage 3 setup screen, after DAM.

Stage 2 (shared setup, per-mode DAM)

Stage 2 uses a single setup screen — one normalization method, one comparison (numerator vs denominator), one DAM method, one FDR method — applied to **both** modes. Inside the orchestrator, two tokio workers run `run_dam` per mode in parallel; the running screen shows two stacked progress bars. If either mode fails, the error message names which mode (Positive: ... or Negative: ...).

The volcano-plot screen renders a POS | NEG tab bar above the plot area. Each tab caches its own texture; changing any threshold slider invalidates both. The PNG export uses a mode-specific default filename (`volcano-pos.png` / `volcano-neg.png`). The DAM CSV export emits a leading `# Mode: dual (POS+NEG)` comment line and prepends a Mode column to every row, with rows ordered POS-first then NEG-second.

Stage 3 — dual-mode N and K math

Stage 3 builds the universe N and foreground K from BOTH modes' DAM features under a conflict-only-strict union rule (the conservative choice: opposing- direction signals exclude a compound).

PubChem and KEGG /conv calls run ONCE on the unioned InChIKey set so the network cost does not double in dual-mode.

N (universe) = union of every cpd reachable from any feature in any mode via the PubChem → KEGG conv chain.

Per-mode trend aggregation. For each cpd *c*, gather the per-feature trends from each mode separately and aggregate:

- Up — any contributing feature in this mode flagged Up, none Down
- Down — symmetric
- NS — only non-significant features
- Conflict — both Up and Down features in the same mode (same-InChIKey-different-trends edge case)
- Absent — the cpd is not reachable from this mode at all

K (foreground) under the conflict-only-strict rule. For active direction Up: a cpd enters K iff at least one mode says Up AND no mode says Down AND no mode is in Conflict. Down is symmetric. Both requires at least one Up or Down signal AND no Conflict AND not (Up in one mode AND Down in another).

Single-mode applies the same conflict rule (since 2026-06-01). A single-mode run is the degenerate one-mode case of this rule: a compound reached by both an Up feature and a Down feature within the single mode — two distinct InChIKeys that map to the **same** KEGG compound, one Up + one Down — aggregates to Conflict and is **excluded** from K, the same conservative choice as dual-mode. (Before this, single-mode kept such ambiguous compounds in K.) The conflict-excluded count appears in the Stage 3 INFO log. Single-mode K is unchanged for any dataset that has no such intra-mode conflict.

The bottom-panel **Data** tab surfaces the dual-mode partition as part of the universe / foreground provenance funnels:

```
Universe — all tested features (measurable metabolome)
... → N KEGG cpds  (POS-only: a; NEG-only: b; in both: c)
Foreground — significant features (active direction)
... → K KEGG cpds  (sig POS-only: d; sig NEG-only: e; agree both: f; excluded by conflict: g)
```

A yellow K source: POS only (NEG had 0 sig features in the active direction) line appears when one mode contributed every K cpd. The enrichment CSV emits a leading `# Mode: dual (POS+NEG)` comment line; the per-row CSV shape is unchanged (the ORA math is mode-agnostic).

Worked example

Using the `data/double-mode/ fixtures` (3 control + 3 treatment biosamples × 2 modes = 12 sample columns + 6 biosamples):

1. Stage 1: load `POS_*.txt` into slot #1 (Mode: Positive), `NEG_*.txt` into slot #2 (Mode: Negative), and `metadata.csv`. Click Continue to DAM.
2. Stage 2: pick treatment vs control, leave normalization and FDR at defaults. The running screen shows two progress bars; expect ~6–60 s per mode depending on feature count.

3. Stage 2 threshold: flip between the POS and NEG tabs to inspect each volcano; download a tabbed PNG or the unified CSV. Click `Continue` to `Enrichment`.
4. Stage 3 setup: pick a KEGG species (Pathway mode) or Level + organism Group (Module mode); the inline progress strip streams the KEGG fetch. After it completes, click `Run Enrichment`.
5. Stage 3 result: the result panel shows the breakdown block; conflict-excluded cpd IDs appear in the log at INFO. Adjust Top N inline if you want fewer/more rows on the dot plot, then click `Re-draw dot plot`. The dot plot keeps the Top N *most significant* entries and stacks them by *fold enrichment* (largest on top); the canvas height re-fits to the rows shown on each redraw (see "Dot plot selection vs ordering" above).

Caches and provenance

The Stage 3 caches (`pubchem.json`, `cid_to_cpd.json`, `modules.json`) store **per-entry** `fetches_at: DateTime<Utc>` timestamps, distinct from the Stage 1 species cache (file-level timestamp). The per-entry granularity is intentional: the caches grow incrementally across many sessions over weeks or months, and a file-level timestamp would either lie about ages or force frequent full-refreshes. The Stage 3 result screen surfaces this as a time span (PubChem: 2026-03-01 -> 2026-05-22); module mode additionally shows the modules cache time span across the **retained** modules used in that run, not the entire cache.

Cache locks:

- **PubChem .inchikey.lock + KEGG .cid_to_cpd.lock** — short-lived, held only during the cache write. 30 s wait with 100 ms polling. (Both files are dot-prefixed / hidden.)
- **KEGG .modules.lock** — long-running advisory lock held for the entire ~6–8 min module fetch. The lock file embeds the holder's PID and a heartbeat `last_seen_at` timestamp rewritten at most every 30 s. A concurrent app instance sees the live lock and waits up to 30 min (5 s polling) for it to clear. If the heartbeat is older than 90 s the lock is treated as orphaned (holder crashed) and overwritten. This prevents two app instances from racing through the module fetch loop and tripping KEGG's 403 rate-limit in tandem.
- **Startup cleanup.** On every application launch, all `.lock` files in the cache directory are removed unconditionally so a crash never permanently blocks future writes.

Cache freshness — **no staleness thresholds**. None of the KEGG caches expire: a cached entry is always returned regardless of age, and the app never silently re-fetches on its own. (The 7-day pathway / 30-day module thresholds present in earlier releases were removed in 2026-05.) Instead the bottom-panel **Data** tab's `Cache data` block (on the `Enrichment Analysis + Enrichment Result` screens) surfaces fetch times neutrally and leaves the refresh decision to you (these fetch-time lines and `Refresh` buttons were relocated out of the screen bodies into the `Data` tab in 2026-05):

- Per-species pathway cache: shows `Cached <ts> (N days ago) (setup) or KEGG pathways (<code>): <ts> (result);` re-fetch via the `Refresh KEGG pathway cache` button.
- Module entries: shows a `KEGG modules fetched date: <oldest> -> <newest>` span; the warm-fetch decision is cache-key membership (is this module already cached?), never age. Re-fetch via the `Refresh KEGG module cache` button.
- On the `Enrichment Result` screen the catalogue-refresh button (module / pathway) navigates back to the `Setup` screen to run the re-fetch there (where its progress strip lives); the PubChem / KEGG-conv refreshes run in place via a confirmation modal.
- Organism list (`organisms.json`): loaded once at startup (cache-first: an on-disk copy always wins regardless of age) with no in-app refresh button. To force a fresh `/list/organism` fetch, delete `organisms.json` from the cache directory and relaunch. (The `Refresh KEGG pathway cache` button refetches only the selected species' pathway→compound map, not the roster.)

Saving and loading session settings (reproducibility)

Two buttons in the `Data` tab — **[Save settings...]** and **[Load settings...]** — let you snapshot every Stage 1–3 parameter to a JSON file and re-apply it later. The intent is reproducibility: if you (or a collaborator) re-run with the same snapshot and the same inputs, the analysis is bit-equal.

What's in the file

A pretty-printed JSON containing:

- `schema_version` (currently 3 — `log_transform: bool` first bumped it 1 → 2 via `add-log-transform-and-scaling`, then `min_entry_size` bumped it 2 → 3 via `add-min-entry-size-filter`, both archived 2026-05-27), `app_version`, `saved_at` (UTC), a `user_note` field initially "" — you can open the file in any text editor and fill it in.
- `input_files` — for each MS-DIAL .txt and the metadata .csv you had loaded at save time: the file's basename + its SHA-256. **Hashes only — your raw data is never included.** This lets a future Load detect when your inputs have drifted from what the snapshot was made against.
- `settings` — every parameter from Stage 1 through Stage 3 (analysis mode, species / organism group, comparison groups, DAM method, normalization, FDR method, thresholds, export sizes, enrichment direction / FDR / top-N).

When is each button available

- **Save settings...** is enabled on every screen after the startup splash, whether or not inputs are loaded. Saving from a blank Stage 1 captures your preferred defaults as a preset for next time.
- **Load settings...** is enabled **only on Stage 1**. On other stages the button is greyed; hovering it shows "Loading settings is only available on the Stage 1 input screen." This is deliberate — applying a snapshot mid-analysis would leave on-screen results out of sync with the new parameters, so the workflow asks you to re-run from inputs.

Loading workflow

1. Click [**Load settings...**] on Stage 1. The OS file picker opens.
2. Pick a saved .json. A confirm modal shows you what's in it:
 - Saved-at timestamp (in your local time), the snapshot's app version, the user note (if any).
 - A one-line summary of the settings (analysis mode, DAM method + FDR, normalization, enrichment direction + FDR + top-N).
 - **Hash mismatches** — if any of your currently-loaded input files have a different SHA-256 from the snapshot's, they're listed here. The settings still apply if you continue, but you're warned that the inputs have drifted.
 - **Field resets** — if the snapshot named a numerator / denominator group, a metadata column, or a PQN reference group that doesn't exist in the metadata you currently have loaded, those fields are listed and reset to None on apply. You'll need to re-pick them at Stage 2 setup. (This section only appears when you have metadata loaded at Load time; if you Load before uploading metadata, the safety net is the Stage 2 setup gate instead — see next paragraph.)
3. Click **Apply settings** to overwrite your current settings, or **Cancel** to discard.

What if I load settings before uploading metadata?

Snapshot's numerator / denominator are written into settings verbatim (no validation happens at Load time because there's no metadata to compare against). When you later upload metadata and advance to Stage 2 setup, the gate checks group membership: if the preserved value doesn't appear in the new metadata's groups, the "Start DAM" button is rendered greyed, with an inline warning ([⚠](#) Numerator/denominator group not present in the loaded metadata.) and the same text as a hover tooltip. Re-pick a valid group from the ComboBox dropdown and the warning clears.

Hand-editing the JSON

The file is plain UTF-8 JSON, pretty-printed. You can:

- Add a comment in the `user_note` field.
- Tweak a single threshold without re-saving from the app.
- Strip the `input_files` block to share a "settings only" snapshot (Load handles an empty `input_files` array — hash check is skipped).

Hand-editing the `schema_version` to a number other than 3, or breaking the JSON syntax, surfaces a clear error toast on Load (e.g. *"This settings file uses schema version 1; this app expects version 3."* or *"Settings file is not valid JSON (line 7 column 15) ..."*). Snapshots saved before `add-min-entry-size-filter` (2026-05-27) carry `schema_version == 1` or `== 2` and will both be rejected — re-save from your current setup to produce a v3 snapshot.

Reporting bugs

If something looks wrong — an unexpected error, a Stage that hangs, results that don't match expectations — the easiest way to get help is to click **[Download bug report...]** in the log pane (bottom of the window, next to the **Clear** button). A confirmation dialog will list the files the resulting zip will contain, then a save-file dialog lets you pick where to put it.

The zip contains exactly eight files:

- `README.txt` — explains the bundle and its privacy boundary.
- `version.txt` — app build info (package version, git SHA, rustc, target).
- `RUST_LOG.txt` — just the `RUST_LOG` directive value, on a single line.
- `KEGG_CACHE_DIR.txt` — just the `KEGG_CACHE_DIR` env value (or `<unset>`). These two are per-variable files (filename = variable name) so no one can mistake the bundle for a full environment dump — only these two named vars are ever included.
- `logs.txt` — every INFO / WARN / ERROR event from this session (HTTP and other low-level dependency chatter is filtered out so the file stays readable).
- `app_state.txt` — which stage you were on and your current settings (analysis mode, species/group, comparison groups, FDR method, thresholds, etc.).
- `input_summary.txt` — the paths and counts of your loaded MS-DIAL files and metadata CSV (paths only — no cell values).
- `cache_summary.txt` — sizes and freshness timestamps of the KEGG / PubChem cache files (no cached content).

Privacy:

- The bundle never includes your raw MS-DIAL .txt input, your metadata CSV, or any prior CSV/PNG exports.
- Absolute paths inside the bundle have your home directory replaced with `~` (e.g. `/Users/alice/Projects/study/POS.txt` becomes `~/Projects/study/POS.txt`) so the bundle does not leak your account/username when shared publicly (GitHub issues, email).
- Only `RUST_LOG` and `KEGG_CACHE_DIR` env vars are surfaced — never the full process environment.

You can safely attach the zip to a GitHub issue or email it without worrying about leaking your experimental data or your machine identity.

Per-session log files are also kept on disk under `<binary-directory>/data/logs/` for 7 days, then auto-deleted at startup. If you want to capture a log from a previous run, look in that directory before reopening the app.

Key references

- **Brunner–Munzel test.** Brunner, E. & Munzel, U. (2000). *The nonparametric Behrens-Fisher problem: Asymptotic theory and a small-sample approximation*. Biometrical Journal 42 (1): 17–25.
- **Cliff's δ .** Cliff, N. (1993). *Dominance statistics: Ordinal analyses to answer ordinal questions*. Psychological Bulletin 114 (3): 494–509.
- **Welch's t-test.** Welch, B. L. (1947). *The generalization of "Student's" problem when several different population variances are involved*. Biometrika 34 (1/2): 28–35.
- **BY FDR.** Benjamini, Y. & Yekutieli, D. (2001). *The control of the false discovery rate in multiple testing under dependency*. Annals of Statistics 29 (4): 1165–1188.
- **BH FDR.** Benjamini, Y. & Hochberg, Y. (1995). *Controlling the false discovery rate: A practical and powerful approach to multiple testing*. Journal of the Royal Statistical Society B 57 (1): 289–300.
- **Hypergeometric ORA.** Khatri, P., Sirota, M. & Butte, A. J. (2012). *Ten years of pathway analysis: Current approaches and outstanding challenges*. PLoS Computational Biology 8 (2): e1002375.
- **KEGG.** Kanehisa, M., Furumichi, M., Sato, Y., Kawashima, M. & Ishiguro-Watanabe, M. (2023). *KEGG for taxonomy-based analysis of pathways and genomes*. Nucleic Acids Research 51 (D1): D587–D592.
- **KEGG MODULE.** Kanehisa, M., Sato, Y., Kawashima, M., Furumichi, M. & Tanabe, M. (2014). *Data, information, knowledge and principle: back to metabolism in KEGG*. Nucleic Acids Research 42 (D1): D199–D205.
- **Quantile normalization.** Bolstad, B. M., Irizarry, R. A., Åstrand, M. & Speed, T. P. (2003). *A comparison of normalization methods for high density oligonucleotide array data based on variance and bias*. Bioinformatics 19 (2): 185–193.

- **Probabilistic Quotient Normalization (PQN).** Dieterle, F., Ross, A., Schlotterbeck, G. & Senn, H. (2006). *Probabilistic quotient normalization as robust method to account for dilution of complex biological mixtures. Application in ¹H NMR metabonomics.* Analytical Chemistry 78 (13): 4281–4290.